


# Actor-Critic

- Actor-Critic (Discrete actions)
- A3C (Asynchronous Advantage Actor-Critic)



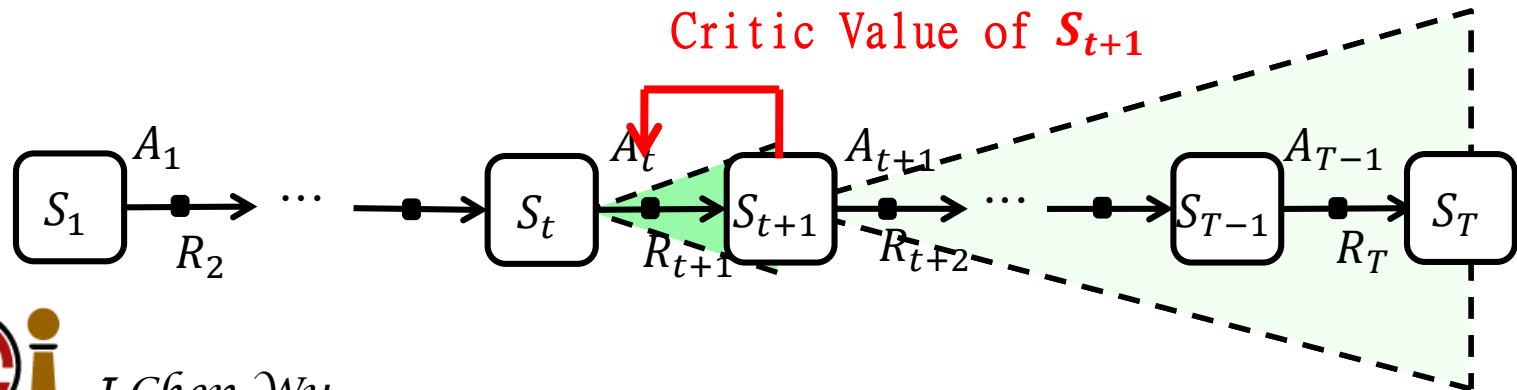
# Problem of REINFORCE

- Problem: Monte-Carlo policy gradient still has **high variance**



- Solution: Actor Critic

- Policy gradient based on the Critic value of  $S_{t+1}$



# Reducing Variance Using a Critic

- We use a **critic** to estimate the action-value function,
$$Q_w(s_t, a_t) \approx Q^{\pi_\theta}(s, a)$$
- Actor-critic algorithms maintain **two** sets of parameters
  - **Critic**: Updates action-value function parameters  $w$
  - **Actor**: Updates policy parameters  $\theta$ , in direction suggested by critic
- Actor-critic algorithms follow **an approximate policy gradient**

$$\nabla_\theta J(\theta) \approx \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) \cdot Q_w(s, a)]$$

$$\Delta\theta = \alpha \nabla_\theta \log \pi_\theta(s, a) \cdot Q_w(s, a)$$



# Estimating the Action-Value Function

- The **critic** is solving a familiar problem: **policy evaluation**
- But, how good is policy  $\pi_\theta$  for current parameters  $\theta$ ?
- This problem was explored in previous two chapters, e.g.
  - Monte-Carlo policy evaluation
  - Temporal-Difference learning
  - TD( $\lambda$ )
- Could also use e.g. least-squares policy evaluation

# Actor-Critic (Discrete Action Space)

- Use two networks: an **actor** and a **critic**

- **Critic** estimates the action-value function

- ▶ Gradient:




$$\nabla_{\omega} L_Q(s_t, a_t | \omega) = ((r_{t+1} + \gamma Q(s_{t+1}, a' | \omega)) - Q(s_t, a_t | \omega)) \nabla_{\omega} Q(s_t, a_t | \omega)$$

- **Actor** updates policy in direction suggested by critic

- ▶ Gradient (approximate policy gradient):

$$J(\theta) = E_{s,a}^{\pi_{\theta}} [Q(s, a | \omega)]$$

$$\nabla_{\theta} J(\theta) = E_{s,a}^{\pi_{\theta}} [\nabla_{\theta} \log \pi(s_t, a_t | \theta) Q(s_t, a_t | \omega)]$$



# Actor-Critic (Discrete Action Space)

- Using linear value function approx.  $Q_w(s, a) = \phi(s, a)^T w$ 
  - **Critic**: Updates  $w$  by linear TD(0)
  - **Actor**: Updates  $\theta$  by policy gradient

**function** QAC

    Initialise  $s, \theta$

    Sample  $a \sim \pi_\theta$

**for** each step **do**

        Sample reward  $r = \mathcal{R}_s^a$ ; sample transition  $s' \sim \mathcal{P}_{s, \cdot}^a$ .

        Sample action  $a' \sim \pi_\theta(s', a')$

$\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$

$\theta = \theta + \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)$

$w \leftarrow w + \beta \delta \phi(s, a)$

$a \leftarrow a', s \leftarrow s'$

**end for**

**end function**



---

# Advantage Actor-Critic



# Reducing Variance Using a Baseline

- Recall:  $\nabla_{\theta} J(\theta) = \mathbb{E}_{s,a}^{\pi_{\theta}} [\nabla_{\theta} \log \pi(s_t, a_t | \theta) Q(s_t, a_t | \omega)]$
- Problem: Can we further reduce variance?
- Solution:
  - This can reduce variance, without changing expectation

$$\begin{aligned} \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) B(s)] &= \sum_{s \in \mathcal{S}} d^{\pi_{\theta}}(s) \sum_a \nabla_{\theta} \pi_{\theta}(s, a) B(s) \\ &= \sum_{s \in \mathcal{S}} d^{\pi_{\theta}}(s) B(s) \nabla_{\theta} \left( \sum_{a \in \mathcal{A}} \pi_{\theta}(s, a) \right) \\ &= 0 \end{aligned}$$

$= 1 \text{ (constant)}$

- Subtract a baseline function  $B(s)$  from the policy gradient
    - A good baseline is the state value function  $B(s) = V^{\pi_{\theta}}(s)$
- So we can rewrite the policy gradient using the **advantage function**  $A^{\pi_{\theta}}(s, a)$

$$\begin{aligned} A^{\pi_{\theta}}(s) &= Q^{\pi_{\theta}}(s, a) - V^{\pi_{\theta}}(s) \\ \nabla_{\theta} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)] \end{aligned}$$



# Estimating the Advantage Function (1)

- The advantage function can significantly reduce variance of policy gradient
- So the critic should really estimate the advantage function
  - For example, by estimating both  $V^{\pi_{\theta}}(s)$  and  $Q^{\pi_{\theta}}(s, a)$
  - Using two function approximators and two parameter vectors,
$$V_v(s) \approx V^{\pi_{\theta}}(s)$$
$$Q_w(s, a) \approx Q^{\pi_{\theta}}(s, a)$$
$$A(s, a) = Q_w(s, a) - V_v(s)$$
- And updating both value functions by e.g. TD learning

## Estimating the Advantage Function (2)

- For the true value function  $V^{\pi_\theta}(s)$ , the TD error  $\delta^{\pi_\theta}$

$$\delta^{\pi_\theta} = r + \gamma V^{\pi_\theta}(s') - V^{\pi_\theta}(s)$$

- is an unbiased estimate of the advantage function

$$\begin{aligned}\mathbb{E}_{\pi_\theta}[\delta^{\pi_\theta} | s, a] &= \mathbb{E}_{\pi_\theta}[r + \gamma V^{\pi_\theta}(s') | s, a] - V^{\pi_\theta}(s) \\ &= Q^{\pi_\theta}(s, a) - V^{\pi_\theta}(s) \\ &= A^{\pi_\theta}(s, a)\end{aligned}$$

- So we can use the TD error to compute the policy gradient

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) \delta^{\pi_\theta}]$$

- In practice we can use an approximate TD error

$$\delta_v = r + \gamma V_v(s') - V_v(s)$$

- This approach only requires one set of critic parameters  $v$



# Critics at Different Time-Scales

- Critic can estimate value function  $V_\theta(s)$  from many targets at different time-scales From last lecture...

- For MC, the target is the return  $v_t$

$$\Delta\theta = \alpha(v_t - V_\theta(s))\phi(s)$$

- For TD(0), the target is the TD target  $r + \gamma V(s')$

$$\Delta\theta = \alpha(r + \gamma V(s') - V_\theta(s))\phi(s)$$

- For forward-view TD( $\lambda$ ), the target is the  $\lambda$ -return  $v_t^\lambda$

$$\Delta\theta = \alpha(v_t^\lambda - V_\theta(s))\phi(s)$$

- For backward-view TD( $\lambda$ ), we use eligibility traces

$$\delta_v = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

$$e_t = \gamma\lambda e_{t-1} + \phi(s_t)$$

$$\Delta\theta = \alpha\delta_t e_t$$



# Actors at Different Time-Scales

- The policy gradient can also be estimated at many time-scales

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

- Monte-Carlo policy gradient uses error from complete return

$$\Delta\theta = \alpha(v_t - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Actor-critic policy gradient uses the one-step TD error

$$\Delta\theta = \alpha(r + \gamma V_v(s_{t+1}) - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Advantage Actor-critic (A2C or A3C) policy gradient uses the  $(k+1)$ -step TD error

$$\Delta\theta = \alpha(v_t^{(k)} - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Some policy gradient algorithms (like PPO) uses TD( $\lambda$ ) error

$$\Delta\theta = \alpha(v_t^{\lambda} - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$



# Actors at Different Time-Scales

- The policy gradient can also be estimated at many time-scales

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

- Monte-Carlo policy gradient uses error from complete return

$$\Delta\theta = \alpha(v_t - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Actor-critic policy gradient uses the one-step TD error

$$\Delta\theta = \alpha(\delta_t) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Advantage Actor-critic (A2C or A3C) policy gradient uses the  $(k+1)$ -step TD error  $= A^{(k+1)}$

$$\Delta\theta = \alpha(\delta_t + \gamma\delta_{t+1} + \cdots + \gamma^k\delta_{t+k}) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Some policy gradient algorithms (like PPO) uses TD( $\lambda$ ) error

$$\Delta\theta = \alpha(\delta_t + \lambda\gamma\delta_{t+1} + \cdots + (\lambda\gamma)^k\delta_{t+k} + \cdots) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

$= A_t^{GAE}$ : Also called GAE (Generalized Advantage Estimator)



# Summary of Policy Gradient Algorithms

- The **policy gradient** has many equivalent forms

|   |                              |
|---|------------------------------|
| $\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) v_t]$ | REINFORCE                    |
| $= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q^w(s, a)]$                     | Q Actor-Critic               |
| $= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \delta]$                        | TD Actor-Critic              |
| $= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{(k)}]$                       | Advantage Actor-Critic       |
| $= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{GAE}]$                       | TD( $\lambda$ ) Actor-Critic |

- Each leads a stochastic gradient ascent algorithm

# Appendix for Advantages and TD( $\lambda$ ) Errors

- TD errors
- n-step TD errors
- GAE
- Eligibility Trace

# Appendix: TD Errors

- TD errors:


$$\delta_t^V = -V(s_t) + r_t + \gamma V(s_{t+1})$$

$$\delta_{t+1}^V = -V(s_{t+1}) + r_{t+1} + \gamma V(s_{t+2})$$

$$\delta_{t+2}^V = -V(s_{t+2}) + r_{t+2} + \gamma V(s_{t+3})$$

$$\vdots$$

$$\delta_{t+n}^V = -V(s_{t+n}) + r_{t+n} + \gamma V(s_{t+n+1})$$





## Appendix: TD Errors

- Weighted TD errors:

$$\delta_t^V = -V(s_t) + r_t + \gamma V(s_{t+1})$$

$$\gamma * \delta_{t+1}^V = \gamma * (-V(s_{t+1}) + r_{t+1} + \gamma V(s_{t+2}))$$

$$\gamma^2 * \delta_{t+2}^V = \gamma^2 * (-V(s_{t+2}) + r_{t+2} + \gamma V(s_{t+3}))$$

$$\vdots$$

$$\gamma^n * \delta_{t+n}^V = \gamma^n * (-V(s_{t+n}) + r_{t+n} + \gamma V(s_{t+n+1}))$$

# Appendix: n-Step TD Errors

- Sum them up, becoming n-step TD errors.

$$\begin{aligned}
 \delta_t^V &= -V(s_t) + r_t + \cancel{\gamma V(s_{t+1})} \\
 \gamma * \delta_{t+1}^V &= \gamma * (\cancel{-V(s_{t+1})} + r_{t+1} + \cancel{\gamma V(s_{t+2})}) \\
 \gamma^2 * \delta_{t+2}^V &= \gamma^2 * (\cancel{-V(s_{t+2})} + r_{t+2} + \cancel{\gamma V(s_{t+3})}) \\
 &\vdots \\
 + \quad \gamma^n * \delta_{t+n}^V &= \gamma^n * (\cancel{-V(s_{t+n})} + r_{t+n} + \cancel{\gamma V(s_{t+n+1})})
 \end{aligned}$$

$$\begin{aligned}
 &\delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V + \dots + \gamma^n \delta_{t+n}^V \\
 &= -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{n+1} V(s_{t+n+1}) \\
 &= -V(s_t) + G_t^{(n+1)} \\
 &= \hat{A}_t^{(n+1)}
 \end{aligned}$$

If  $n \rightarrow \infty$ , it becomes MC learning. Why?



## Appendix: n-Step TD Errors

- n-step TD errors:

$$\begin{aligned}\hat{A}_t^{(1)} &:= \delta_t^V \\ \hat{A}_t^{(2)} &:= (\delta_t^V + \gamma \delta_{t+1}^V) \\ \hat{A}_t^{(3)} &:= (\delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V) \\ &\vdots \\ \hat{A}_t^{(n)} &:= \sum_{k=1}^n \gamma^{k-1} * \delta_{t+k-1}^V\end{aligned}$$

# Appendix: n-Step TD Errors and GAE

- Weighted n-step TD errors:
  - The same trick as TD( $\lambda$ )
- Then, sum them up.

$$\begin{aligned}
 (1 - \lambda) * \hat{A}_t^{(1)} &:= (1 - \lambda) & * \delta_t^V \\
 (1 - \lambda)\lambda * \hat{A}_t^{(2)} &:= (1 - \lambda)\lambda & * (\delta_t^V + \gamma\delta_{t+1}^V) \\
 (1 - \lambda)\lambda^2 * \hat{A}_t^{(3)} &:= (1 - \lambda)\lambda^2 & * (\delta_t^V + \gamma\delta_{t+1}^V + \gamma^2\delta_{t+2}^V)
 \end{aligned}$$

$$\vdots$$

$$+ \quad (1 - \lambda)\lambda^{n-1} * \hat{A}_t^{(n)} := (1 - \lambda) \sum_{k=1}^n \gamma^{k-1} * \delta_{t+k-1}^V$$

$$\hat{A}_t^{GAE(\gamma, \lambda)} = (1 - \lambda)(\hat{A}_t^{(1)} + \lambda\hat{A}_t^{(2)} + \lambda^2\hat{A}_t^{(3)} + \dots + \lambda^{n-1}\hat{A}_t^{(n)} + \dots)$$



# Appendix: n-Step TD Errors and GAE

- The sum of exponentially-weighted TD residuals denoted as  $\hat{A}_t^{GAE(\gamma, \lambda)}$  (actually equals to  $G_t^\lambda - V(S_t)$  for  $TD(\lambda)$ )

$$\begin{aligned}
 \hat{A}_t^{GAE(\gamma, \lambda)} &= (1 - \lambda) \left( \hat{A}_t^{(1)} + \lambda \hat{A}_t^{(2)} + \lambda^2 \hat{A}_t^{(3)} + \dots + \lambda^{n-1} \hat{A}_t^{(n)} + \dots \right) \\
 &= (1 - \lambda) \left( (\delta_t^V) + \lambda (\delta_t^V + \gamma \delta_{t+1}^V) + \lambda^2 (\delta_t^V + \gamma \delta_{t+1}^V + \gamma \delta_{t+2}^V) + \dots \right) \\
 &= (1 - \lambda) \left( \begin{array}{c} \delta_t^V (1 + \lambda + \lambda^2 + \dots) + \\ \gamma \lambda \delta_{t+1}^V (1 + \lambda + \lambda^2 + \dots) + \\ (\gamma \lambda)^2 \delta_{t+2}^V (1 + \lambda + \lambda^2 + \dots) + \\ \dots \end{array} \right) \\
 &= (1 - \lambda) \left( \delta_t^V \left( \frac{1}{1 - \lambda} \right) + \gamma \lambda \delta_{t+1}^V \left( \frac{1}{1 - \lambda} \right) + (\gamma \lambda)^2 \delta_{t+2}^V \left( \frac{1}{1 - \lambda} \right) + \dots \right) \\
 &= \sum_{n=0}^{\infty} (\gamma \lambda)^n \delta_{t+n}^V = \delta_t^V + \lambda \gamma \delta_{t+1}^V + \dots + (\lambda \gamma)^k \delta_{t+k}^V + \dots
 \end{aligned}$$



# Appendix: Recall $TD(\lambda)$

- $\lambda$ -return  $G_t^\lambda$ :
  - combines all  $n$ -step returns  $G_t^{(n)}$
- Using weight  $(1 - \lambda) \lambda^{n-1}$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$


or (in the case of termination)

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + (1 - \lambda) \sum_{n=T-t-1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t-1} G_t$$

- Forward-view  $TD(\lambda)$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$




# Appendix: GAE and Eligibility Trace

- Eligibility trace:

$$E_0(s) = 0$$

$$E_t(s) = (\gamma \lambda) E_{t-1}(s) + 1(S_t = s)$$



$$\begin{aligned} \hat{A}_t^{GAE(\gamma, \lambda)} &= 1 \delta_t^V + \lambda \gamma \delta_{t+1}^V + (\lambda \gamma)^2 \delta_{t+2}^V + \dots + (\lambda \gamma)^k \delta_{t+k}^V + \dots \\ \hat{A}_{t+1}^{GAE(\gamma, \lambda)} &= 1 \delta_{t+1}^V + \lambda \gamma \delta_{t+2}^V + \dots + (\lambda \gamma)^{k-1} \delta_{t+k}^V + \dots \\ \hat{A}_{t+2}^{GAE(\gamma, \lambda)} &= 1 \delta_{t+2}^V + \dots + (\lambda \gamma)^{k-2} \delta_{t+k}^V + \dots \\ &\dots \end{aligned}$$

$E_t(s_t)$        $E_t(s_{t+1})$        $E_t(s_{t+2})$

---


# A3C (Asynchronous Advantage Actor-Critic)






# Asynchronous Advantage Actor-critic (A3C)

- Asynchronous Lock-Free Reinforcement Learning
  - Use two main ideas to make the algorithm practical:
    - ▶ Multiple threads on a **single machine**
    - ▶ Multiple actor-learners applying **online updates** in parallel (**no experience replay**)



# Asynchronous Advantage Actor-critic (A3C)

- Instead of experience replay, we **asynchronously** execute **multiple agents** in parallel.
  - Decorrelate the agents' data into a more stationary process
  - Enable a much larger spectrum of fundamental **on-policy RL algorithms**
- For each worker (asynchronous part):
  - Copy all parameters from the global network.
  - keep playing and computing gradients.
  - Every N iterations:
    1. **Update** all gradients to the global network.
    2. **Copy** all new parameters from the global network



# Asynchronous Advantage Actor-critic (A3C)

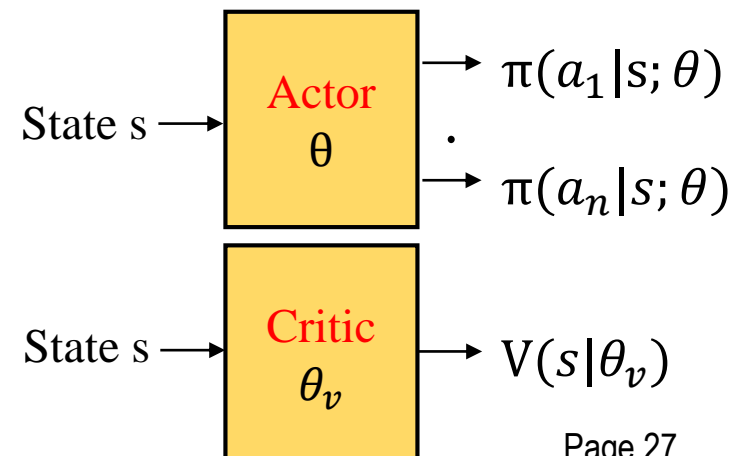
- Asynchronous advantage actor-critic(A3C) maintains a policy  $\pi(a_t|s_t; \theta)$  and an estimate of the value function  $V(s_t, \theta_v)$ .

- The update performed by the algorithm can be seen as

$$\nabla_{\theta} \log \pi(a_t|s_t; \theta) \underbrace{A(s_t, a_t; \theta, \theta_v)}_{\text{Advantage}} = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta_v) - V(s_t; \theta_v)$$

– Make  $k$ -step operations, and then calculate advantages backwards.

- Intuitively, the network should be




# Asynchronous Advantage Actor-critic (A3C)

- Asynchronous advantage actor-critic(A3C) maintains a policy  $\pi(a_t|s_t; \theta)$  and an estimate of the value function  $V(s_t, \theta_v)$ .

- The update performed by the algorithm can be seen as

$$\nabla_{\theta} \log \pi(a_t|s_t; \theta) \underbrace{A(s_t, a_t; \theta, \theta_v)}_{\text{Advantage}} = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta_v) - V(s_t; \theta_v)$$

- Make  $k$ -step operations, and then calculate advantages backwards.
- Typically use a convolutional neural network that has two heads:
  - one **softmax output** for the policy  $\pi(a_t|s_t; \theta)$
  - one **output** for the value function  $V(s_t; \theta_v)$
  - **all non-output layers are shared**



# Asynchronous Advantage Actor-critic (A3C)

**repeat**

Sync  $\theta' = \theta, \theta'_v = \theta_v$

$t_{start} = t$

Get state  $s_t$

**repeat** (note:  $t = t + 1$ )

Perform  $a_t$  according to policy  $\pi(a_t|s_t; \theta')$

Receive  $s'$  and reward  $r$

**until** terminal  $s_t$  or  $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s' \\ V(s_t, \theta'_v) & \text{for non-terminal } s' \end{cases}$

(note:  $t = t_{start} + t_{max}$ ,  
if not terminal)

**for**  $i \in \{t - 1, \dots, t_{start}\}$  **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt  $\theta'$ :  $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta') (R - V(s_i; \theta'_v))$

Accumulate gradients wrt  $\theta'_v$ :  $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

**end for**

Perform asynchronous update of  $\theta$  using  $d\theta$  and of  $\theta_v$  using  $d\theta_v$ .

**until**  $T \rightarrow T_{max}$

$\theta, \theta_v$ : global shared parameters

$T$ : global shared counter

$\theta', \theta'_v$ : thread specific parameters

$t$ : thread step counter



# Experiments – A3C

| Method                                 | Training Time        | Mean   | Median |
|--|----------------------|--------|--------|
| DQN (from [Nair et al., 2015])         | 8 days on GPU        | 121.9% | 47.5%  |
| Gorila [Nair et al., 2015]             | 4 days, 100 machines | 215.2% | 71.3%  |
| Double DQN [Van Hasselt et al., 2015]  | 8 days on GPU        | 332.9% | 110.9% |
| Dueling Double DQN [Wang et al., 2015] | 8 days on GPU        | 343.8% | 117.1% |
| Prioritized DQN [Schaul et al., 2015]  | 8 days on GPU        | 463.6% | 127.6% |
| A3C, FF                                | 1 day on CPU         | 344.1% | 68.2%  |
| A3C, FF                                | 4 days on CPU        | 496.8% | 116.6% |
| A3C, LSTM                              | 4 days on CPU        | 623.0% | 112.6% |

Table 1: Mean and median human-normalized scores on 57 Atari games using the human starts evaluation metric.

