

Reinforcement Learning with Demonstration

Acknowledgement: Most slides were contributed by
施園維、蔡承倫、郭奎廷、林九州、吳岱霖 etc.,
and organized by 廖唯辰.



Outline

- DQfD
- R2D3
- Decision Transformer

Reference

- “Deep Q-learning from Demonstrations”
 - Published at AAAI 2018
 - Todd Hester, Matej Vecerik, Olivier Pietquin, Marc Lanctot, Tom Schaul, Bilal Piot, Dan Horgan, John Quan, Andrew Sendonaris, Gabriel Dulac-Arnold, Ian Osband, John Agapiou, Joel Z. Leibo, Audrunas Gruslys
 - Provided by Google DeepMind
- Slides by Yu-Wei Shih

Reference

- “Making Efficient Use of Demonstrations to Solve Hard Exploration Problems.”
 - Published at ICLR 2020
 - Gulcehre, Caglar, Tom Le Paine, Bobak Shahriari, Misha Denil, Matt Hoffman, Hubert Soyer, Richard Tanburn, Steven Kapturowski, Neil Rabinowitz, and Duncan Williams
 - Provided by Google DeepMind
- Slides by 郭奎廷

Reference

- “Decision Transformer: Reinforcement Learning via Sequence Modeling.”
 - Published at NeurIPS 2021
 - Lili Chen, Kevin Lu, Aravind Rajeswaran, Kimin Lee, etc.
 - UC Berkeley, Facebook AI Research, UCLA, OpenAI, Google Brain
- Slides by Tai-Lin Wu



DQfD: Deep Q-learning from Demonstrations

Introduction

- DRL agents usually have poor performance during the early stage of training
 - This could be problematic for many real world tasks that require fine performance at the beginning
- DQfD leverages small demonstration datasets to greatly speed up early phase training process

Baseline: PDD DQN

- Prioritized Experience Replay + Dueling Network + DDQN
- Based on this structure, DQfD makes several improvements to leverage demonstration data

Note:

Dueling Network is used in experiment but the authors didn't mention this as a requirement in DQfD.


DQfD: Four Losses

- DQfD uses 4 losses to update:
 - 1-step double Q-learning loss
 - n-step double Q-learning loss
 - Supervised large margin classification loss
 - L2 regularization loss




1-step Double Q-learning Loss

Q (behavior network)



Q' (target network)




$$y_t = R(s_t, a_t) + \gamma Q'(s_{t+1}, \argmax_a Q(s_{t+1}, a))$$

$$J_{DQ}(Q) = \left(y_t - Q(s_t, a_t) \right)^2$$


- Same as the loss in DDQN method.

n-step Double Q-learning Loss

Q (behavior network)



Q' (target network)



$$y_t = R_t + \gamma R_{t+1} + \dots + \gamma^{n-1} R_{t+n-1} + \gamma^n Q'(s_{t+n}, \operatorname{argmax}_a Q(s_{t+n}, a))$$

$$\text{where } R_t = R(s_t, a_t)$$

$$J_n(Q) = (y_t - Q(s_t, a_t))^2$$

- Adding n-step returns helps propagate the values of the expert's trajectory to all the earlier states, leading to better pre-training.
- This paper uses 10-step loss ($n=10$)




Supervised Large Margin Classification Loss

- Purpose: make the Q value of a_E at least a margin larger than Q value of any other actions.
 - a_E is demonstration data
- Encourage the agent to follow demonstration data instead of choosing other values.

$$J_E(Q) = \max_{a \in A} [Q(s, a) + l(a_E, a)] - Q(s, a_E)$$

$$l(a_E, a) = \begin{cases} 0 & \text{if } a = a_E \\ \text{positive value} & \text{if } a \neq a_E \end{cases}$$



Supervised Large Margin Classification Loss


$$Q(s, a_E) \quad \text{1.0}$$

$$Q(s, a) \quad \text{2.8}$$

$$Q(s, a') \quad \text{1.5}$$

$$J_E(Q) = \max_{a \in A} [Q(s, a) + l(a_E, a)] - Q(s, a_E) > 0$$

$$l(a_E, a) = \begin{cases} 0 & \text{if } a = a_E \\ 0.8 & \text{if } a \neq a_E \end{cases}$$



Supervised Large Margin Classification Loss

$$Q(s, a_E) \text{ (blue bar)} = 1.2$$


$$Q(s, a) \text{ (blue bar)} + 0.6 \text{ (orange bar)} = 1.8$$

$$Q(s, a') \text{ (blue bar)} + 0.6 \text{ (orange bar)} = 1.5$$


$$J_E(Q) = \max_{a \in A} [Q(s, a) + l(a_E, a)] - Q(s, a_E) > 0$$

1.8 1.2

$$l(a_E, a) = \begin{cases} 0 & \text{if } a = a_E \\ 0.8 & \text{if } a \neq a_E \end{cases}$$




Supervised Large Margin Classification Loss



$$J_E(Q) = \max_{a \in A} [Q(s, a) + l(a_E, a)] - Q(s, a_E) = 0$$

The red box highlights the term $\max_{a \in A} [Q(s, a) + l(a_E, a)]$ with a value of 1.8. The blue box highlights $Q(s, a_E)$ with a value of 1.8.

$$l(a_E, a) = \begin{cases} 0 & \text{if } a = a_E \\ 0.8 & \text{if } a \neq a_E \end{cases}$$



L2 Regularization Loss

- Use L2 Norm of all weights and bias in model as a regularization loss to avoid over-fitting.



DQfD: Four Losses

- Combining all four losses:

$$J(Q) = J_{DQ}(Q) + \lambda_1 J_n(Q) + \lambda_2 J_E(Q) + \lambda_3 J_{L2}(Q)$$

When updating by demonstration data.

In this paper: $\lambda_1 = 1, \lambda_2 = 1, \lambda_3 = 10^{-5}$

DQfD: Four Losses

- Combining all four losses:

$$J(Q) = \underbrace{J_{DQ}(Q) + \lambda_1 J_n(Q)}_{\text{When updating by self-generated data.}} + \lambda_2 J_E(Q) + \underbrace{\lambda_3 J_{L2}(Q)}_{\text{When updating by self-generated data.}}$$

When updating by self-generated data.

In this paper:

$$\lambda_1 = 1, \lambda_2 = 0, \lambda_3 = 10^{-5}$$



Prioritized Experience Replay (PER)

- Sample important trajectories more frequently.
- Usually importance = |TD error|.

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}, \quad p_i = |\delta_i| + \epsilon$$

$$w_i = \left(\frac{1}{N} \times \frac{1}{P(i)} \right)^\beta$$

P(i) = probability to be sampled for trajectory i

δ = TD error

N = size of replay buffer

α = priority exponent

β = importance sampling exponent

Prioritized Experience Replay (PER)

- Sample important trajectories more frequently.
- Usually importance = |TD error|.

use different constant $\epsilon_a \epsilon_d$ for self-play data and demonstration data ($\epsilon_a = 0.001, \epsilon_d = 1$)

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}, \quad p_i = |\delta_i| + \epsilon$$

$$w_i = \left(\frac{1}{N} \times \frac{1}{P(i)} \right)^\beta$$

$P(i)$ = probability to be sampled for trajectory i

δ = TD error

N = size of replay buffer

α = priority exponent

β = importance sampling exponent

Pseudocode

$k=750,000$ in this paper

- Demonstration data are placed in replay buffer at the beginning of training.
- When replay buffer is full, replace old self-play data but never remove demonstration data.

Pretraining Phase

$$J(Q) = J_{DQ}(Q) + \lambda_1 J_n(Q) + \lambda_2 J_E(Q) + \lambda_3 J_{L2}(Q)$$

Note: in this paper: $k = 750000$

Interacting with
environment

Algorithm 1 Deep Q-learning from Demonstrations.

```

1: Inputs:  $\mathcal{D}^{replay}$ : initialized with demonstration data set,
    $\theta$ : weights for initial behavior network (random),  $\theta'$ :
   weights for target network (random),  $\tau$ : frequency at
   which to update target net,  $k$ : number of pre-training
   gradient updates
2: for steps  $t \in \{1, 2, \dots, k\}$  do
3:   Sample a mini-batch of  $n$  transitions from  $\mathcal{D}^{replay}$ 
   with prioritization
4:   Calculate loss  $J(Q)$  using target network
5:   Perform a gradient descent step to update  $\theta$ 
6:   if  $t \bmod \tau = 0$  then  $\theta' \leftarrow \theta$  end if
7: end for
8: for steps  $t \in \{1, 2, \dots\}$  do
9:   Sample action from behavior policy  $a \sim \pi^{Q_\theta}$ 
10:  Play action  $a$  and observe  $(s', r)$ .
11:  Store  $(s, a, r, s')$  into  $\mathcal{D}^{replay}$ , overwriting oldest
   self-generated transition if over capacity
12:  Sample a mini-batch of  $n$  transitions from  $\mathcal{D}^{replay}$ 
   with prioritization
13:  Calculate loss  $J(Q)$  using target network
14:  Perform a gradient descent step to update  $\theta$ 
15:  if  $t \bmod \tau = 0$  then  $\theta' \leftarrow \theta$  end if
16:   $s \leftarrow s'$ 
17: end for
  
```





Compare with Baseline PDD DQN

- Full DQfD with human demonstration data
- PDD DQN without any demonstration data
- Supervised imitation learning from demonstration data without any interaction with environments.



Result




Demonstration Data Ratio in a Mini-batch



- In environments with sparse rewards, **the constants ϵ in PER** become specifically important.
- Big constant for demonstration data help DQfD agents overcome environments with such rewards by following demonstration data.

$$p_i = |\delta_i| + \epsilon$$

($\epsilon_a = 0.001$, $\epsilon_d = 1$)



Ablation Study



Conclusion

- DQfD has good performance from the beginning of the training due to pretraining by demonstration data
- Adjusting the ratio between demonstration data and self-play data in each mini-batch by Prioritized Experience Reply
- Combining supervised loss and TD loss to follow demonstrators' actions

R2D3

Recurrent Replay Distributed DQN from
Demonstration

(R2D2: Recurrent Replay Distributed DQN)



Introduction

- Sparse reward problem is a challenge for RL methods
 - Common approaches
 - ▶ Intrinsic motivation, reward shaping, or curriculum tasks
 - Drawbacks:
 - ▶ Do not scale well
 - ▶ Lead to unexpected behavior
 - ▶ Difficult to specify
- Learning from demonstrations has proven to be an effective strategy

Introduction

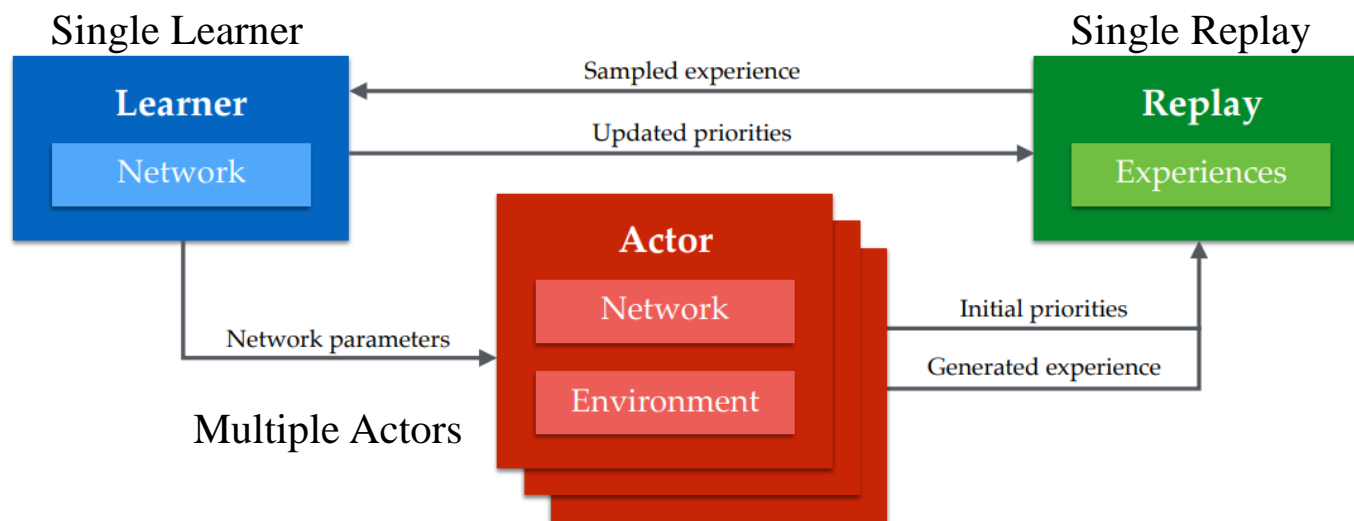
- Three aspects to make learning from demonstrations challenging
 - Sparse rewards
 - Partial observability
 - Highly variable initial conditions
- The demonstrations cannot account for all possible configurations

Background – DQfD

- Leverage expert demonstration datasets to speed up training
- Encourage the agent to follow demonstration data
 - Use supervised imitation loss in training
- Four losses (weighted sum):
 - 1-step double Q-learning loss
 - n-step double Q-learning loss
 - Supervised large margin classification loss
 - L2 regularization loss

Background – Ape-X

- Leverage prioritized experience replay (PER) on distributed architecture




Background – Ape-X

● Problem

- Different games require a significantly different degree of exploration

● Solution

- Give the different actors different exploration policies
 - ▶ E.g., different ϵ values for ϵ -greedy
- Sample with priority to get the most useful data



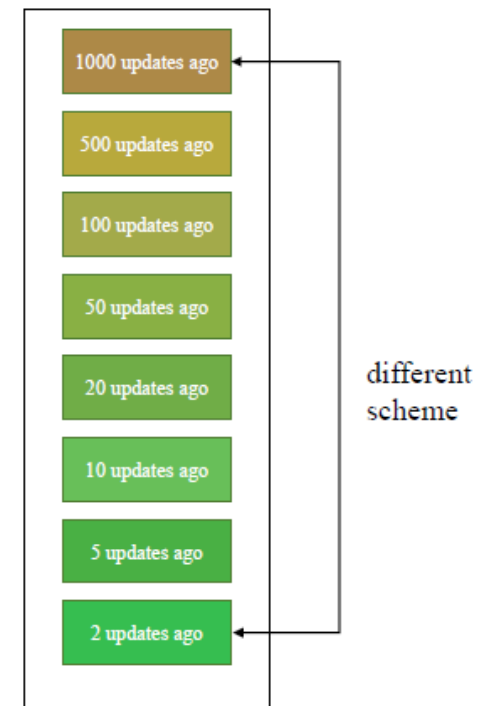
Background – R2D2

- R2D2: Recurrent Replay Distributed DQN
- Build on Ape-X
- Train an RNN from experience replay
 - Partial observability
 - Better representation learning

Training Recurrent RL Agent with Experience Replay

● Stored state

- Storing the recurrent state in replay
 - ▶ Use it to initialize the network at training time
- New problems
 - ▶ Representational drift
 - Out of date representation



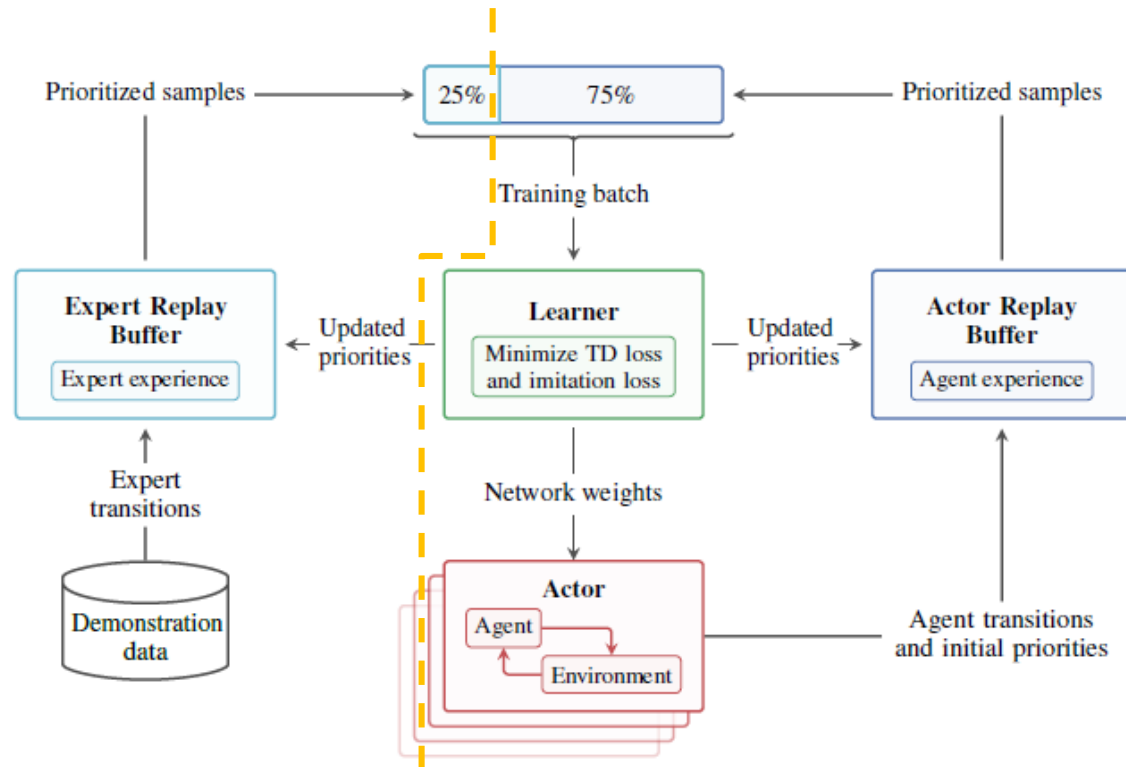
Training Recurrent RL Agent with Experience Replay

● Burn in

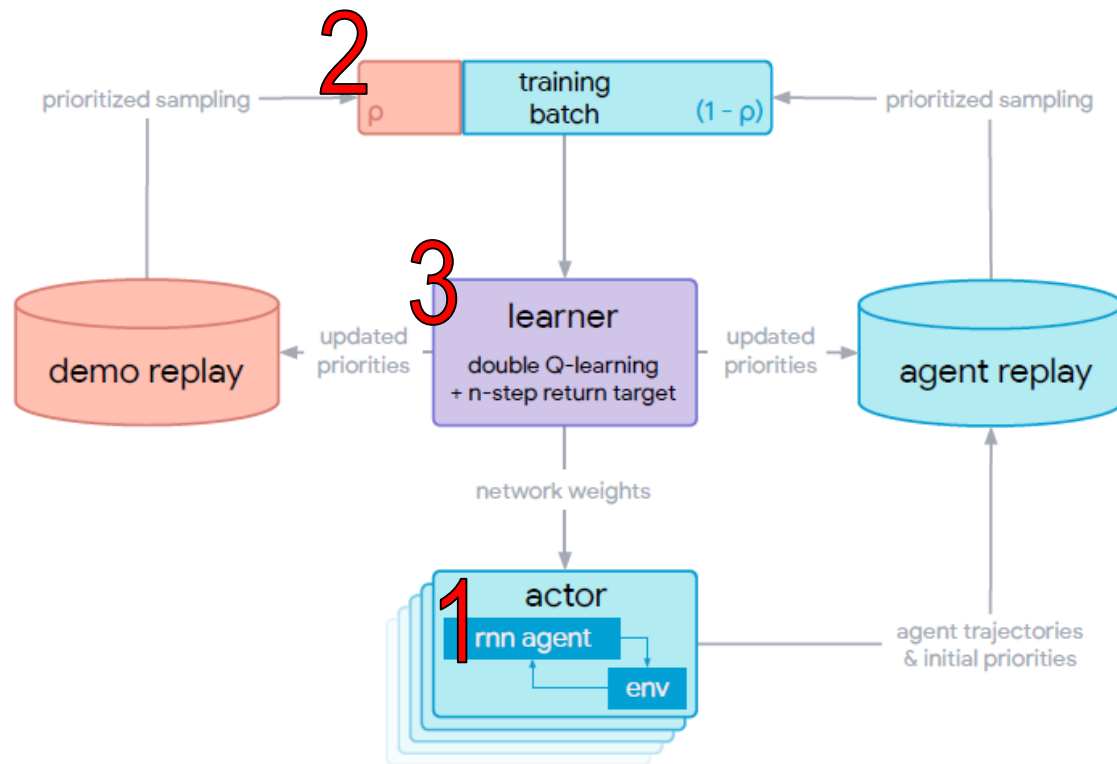
- Allow the network a ‘burn in period’ by using a portion of the replay sequence
- Allow the network to partially recover from a poor start state



Background – Ape-X DQfD



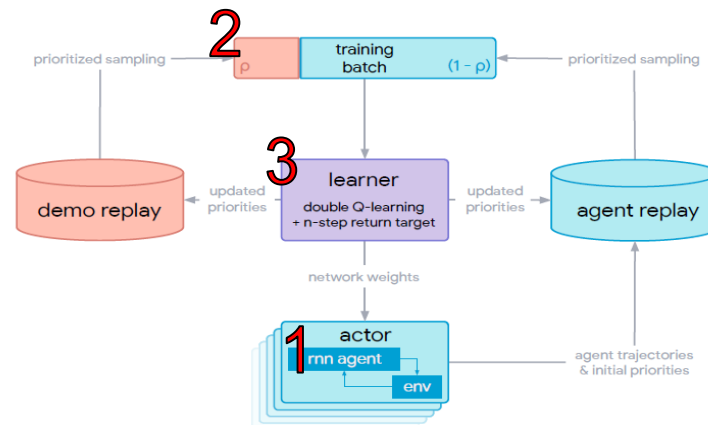
R2D3



R2D3


● Three main changes

1. Use R2D2 to replace Ape-X
 - Deal with partial observable environment
2. Tune the demo ratio ρ
 - Surprisingly, find that the optimal demo ratio is very small (1/256)
3. Remove imitation loss
 - Use expert demos to bias the agent's own autonomous exploration




R2D3 – Hard-Eight Task Suite

- A procedurally-generated 3D world
- Common properties
 - Sparse reward
 - Partially observable
 - Highly variable initial conditions



baseball




drawbridge




remember sensor




throw across




navigate cubes



push blocks



wall sensor




wall sensor stack

R2D3 – Hard-Eight Task Suite

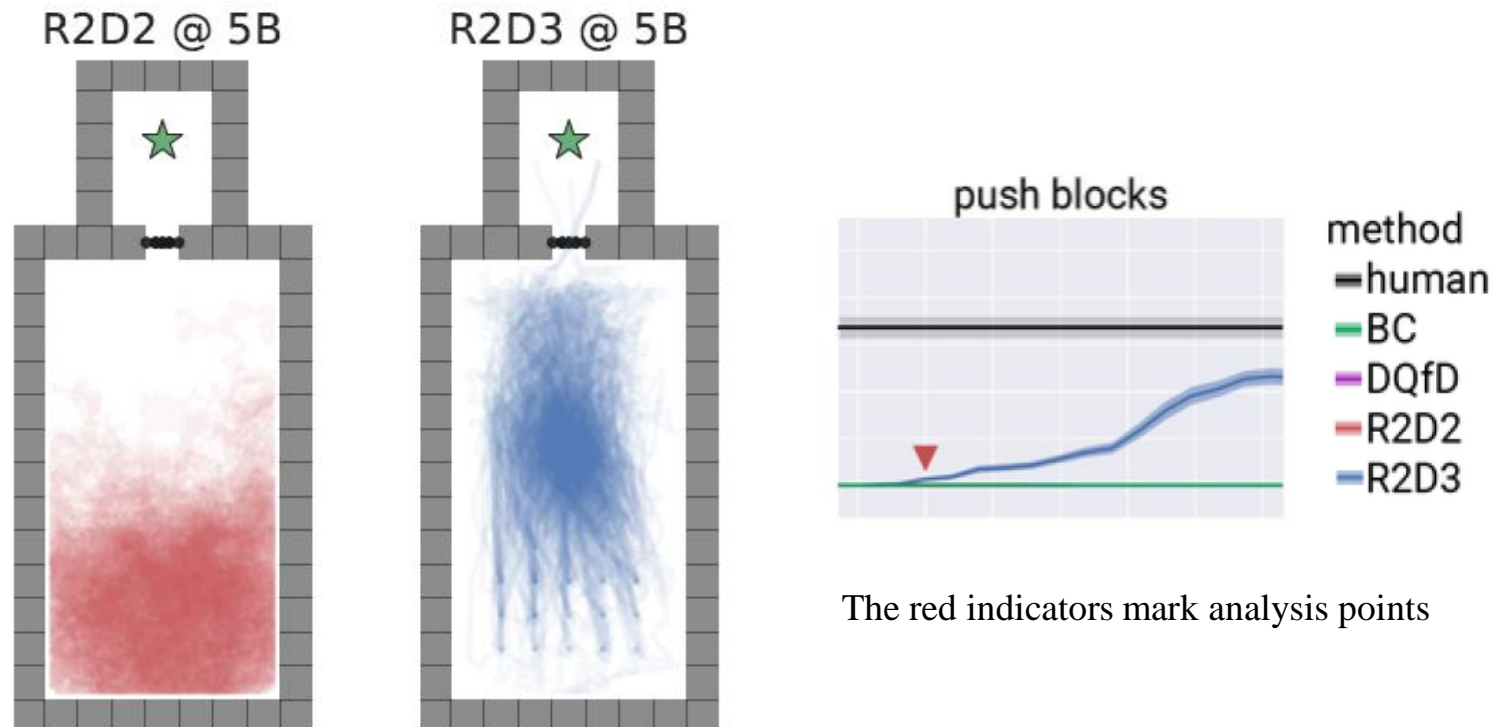
- Example: Push Blocks
- Step
 1. Check the color of sensor
 2. Push a block whose color matches the sensor into the recess
 3. Collect the apple




Experiment – R2D3 Performance in Hard-Eight suite



Experiment – R2D3 Guided Exploration Behavior



Experiment – R2D3 Demo Ratio Tuning



Conclusion


- Design to make efficient use of demonstrations to learn
 - Partially observable environments
 - Sparse rewards
 - Highly variable initial conditions
- Use the expert demonstrations in a way that guides the agent's own autonomous exploration

Decision Transformer

Decision Transformer: Reinforcement Learning
via Sequence Modeling


Introduction

- Recent work has shown transformers can model distributions of **semantic concepts**, and associated advances in language modeling such as GPT-x and BERT have achieved **impressive results**. (GPT3/3.5)




Introduction

- Recent work has shown transformers can model distributions of **semantic concepts**, and associated advances in language modeling such as GPT-x and BERT have achieved **impressive results**. (GPT4)




Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).




Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



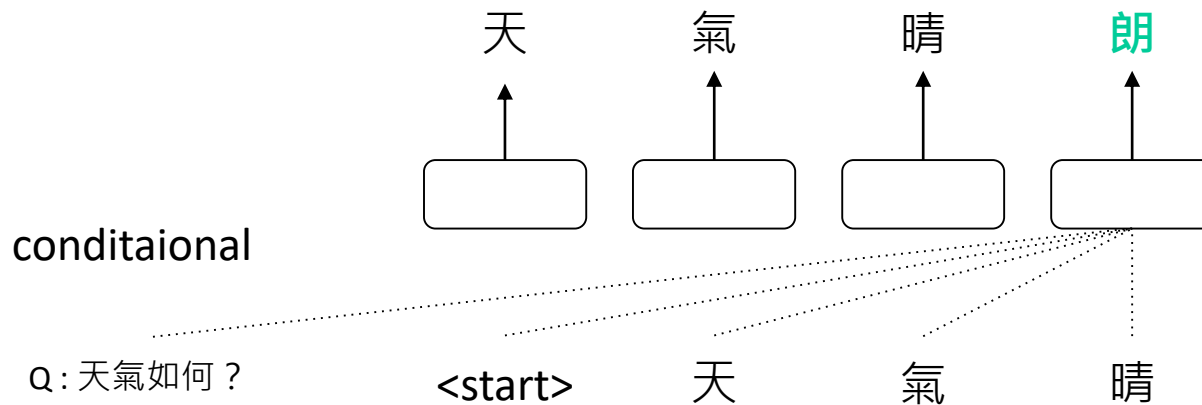
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



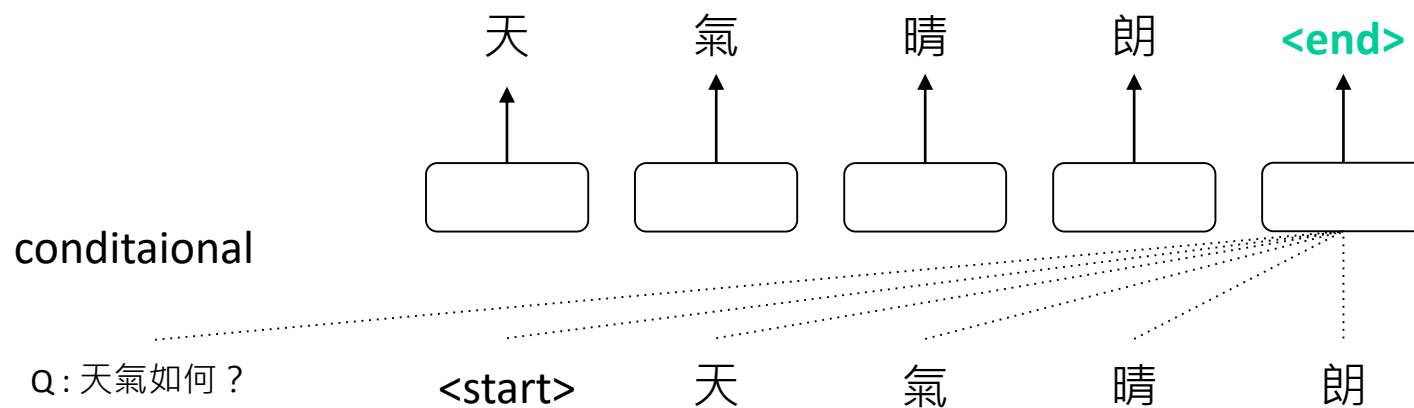
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



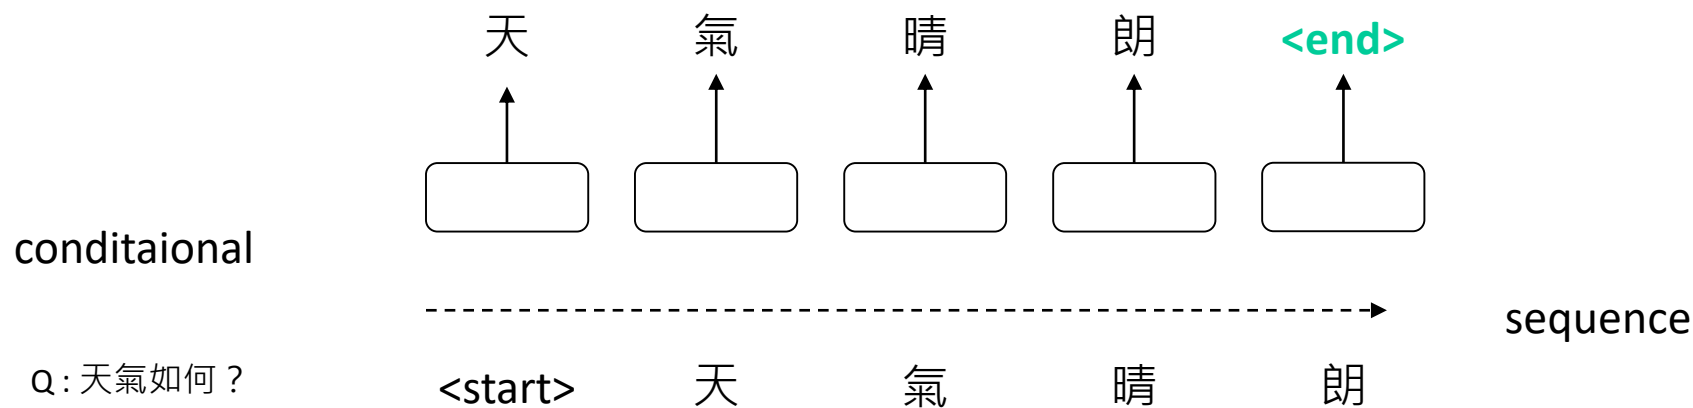
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



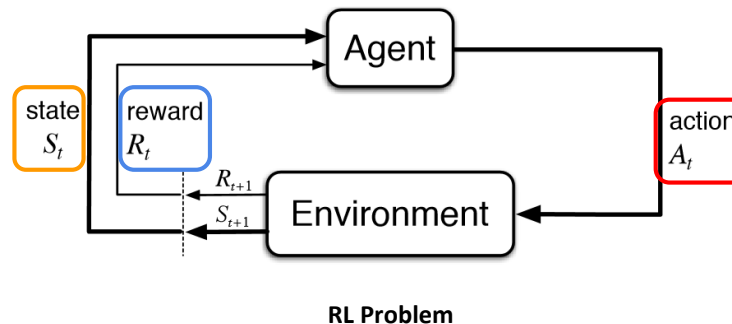
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



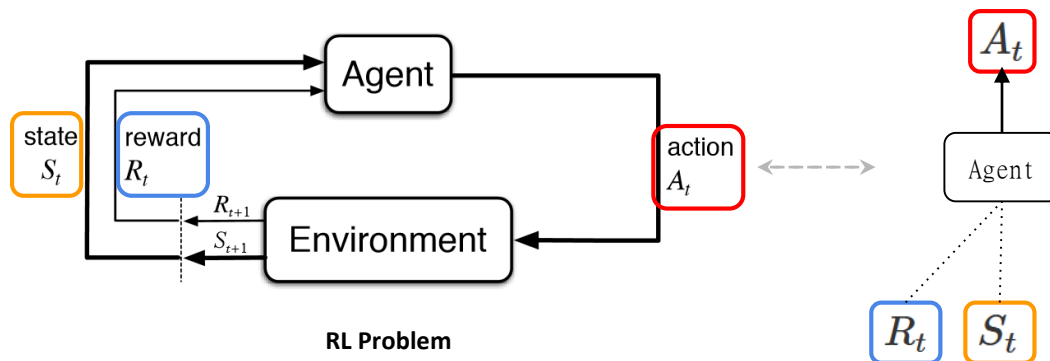
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



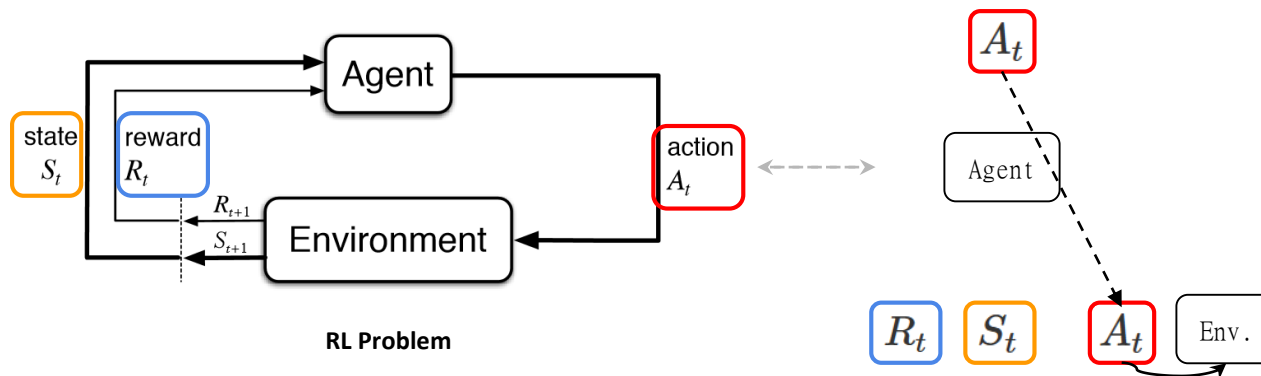
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



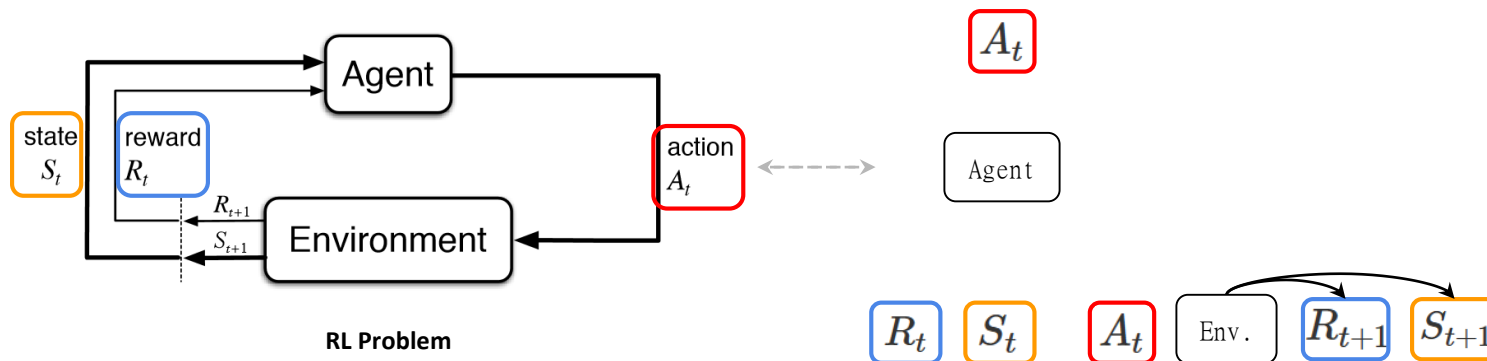
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



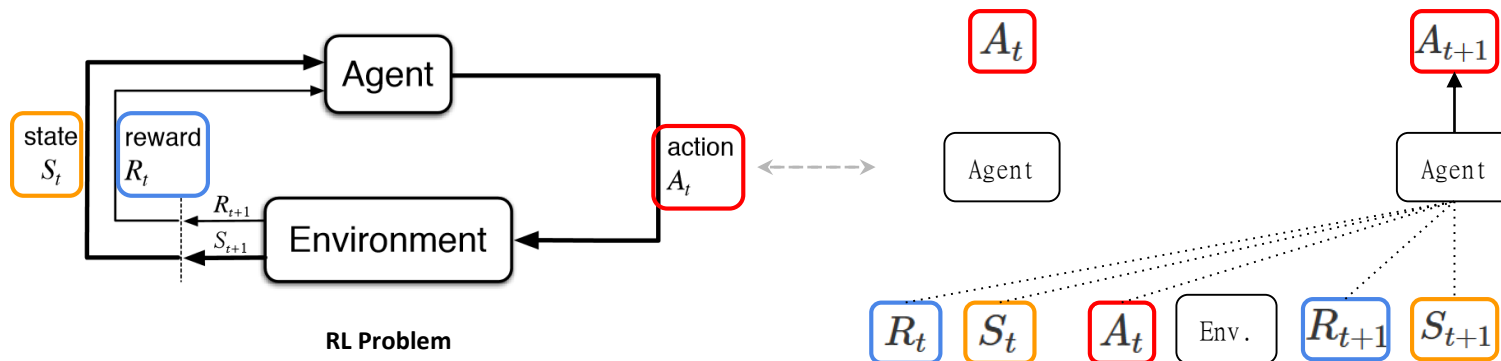
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



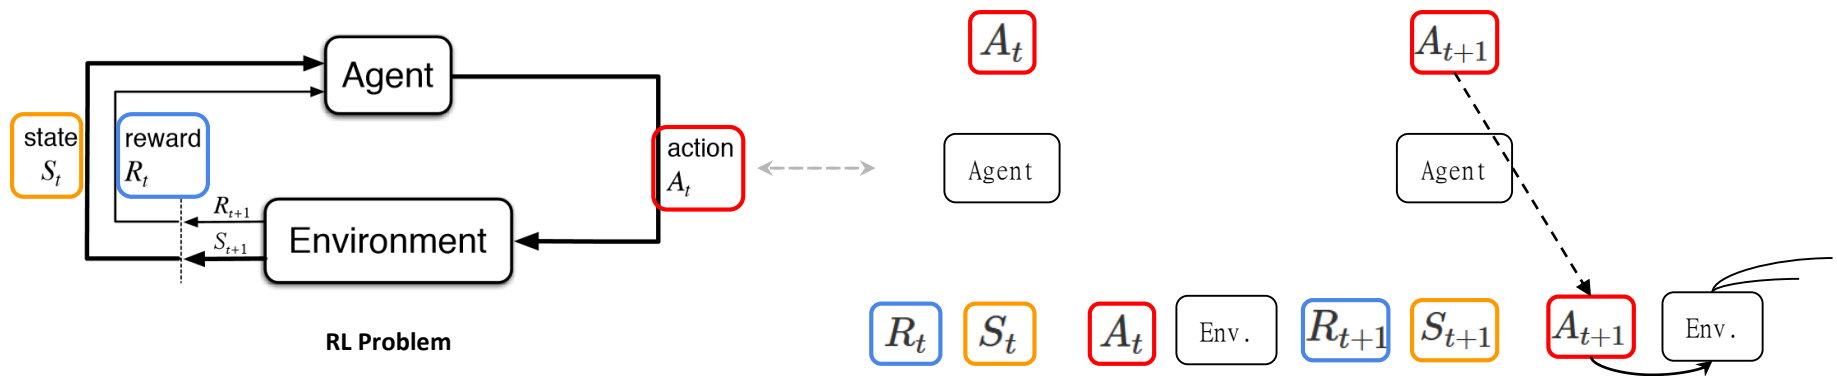
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



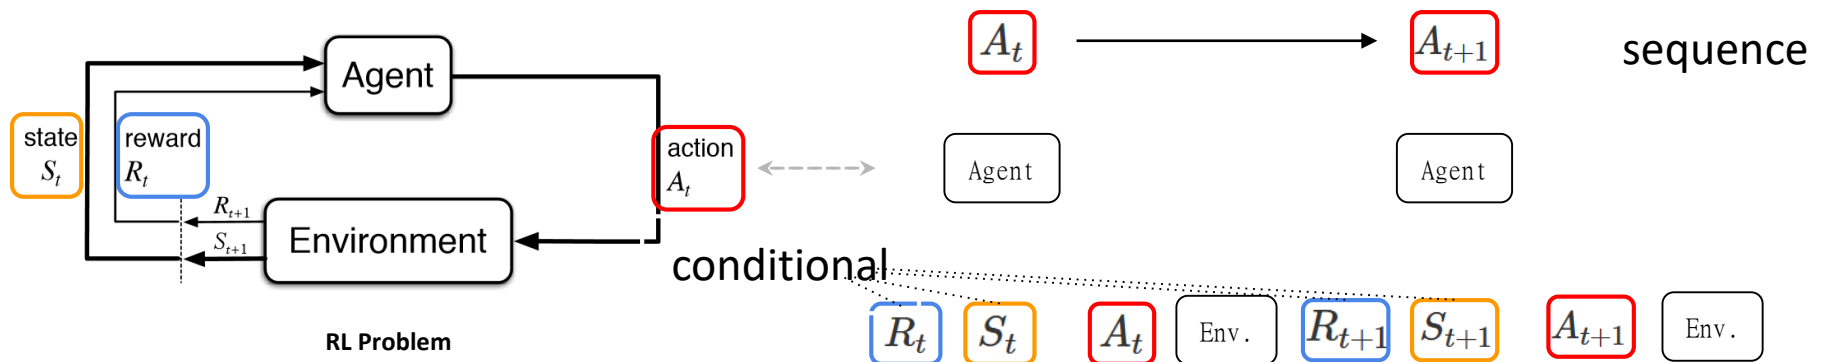
Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).



Introduction

- This paper present Decision Transformer as:
 - An architecture that casts the problem of RL as **conditional sequence modeling** (Generative Transformer).




Introduction

- Decision Transformer **matches** or **exceeds** the performance of state-of-the-art **model-free offline RL baselines** on Atari (2021), OpenAI Gym.




Motivation

- **Sparse rewards:** which refers to infrequent or insufficient reward signals that make it difficult to guide learning
- Common methods used to address the issue:
 - Reward shaping, curriculum learning, etc.



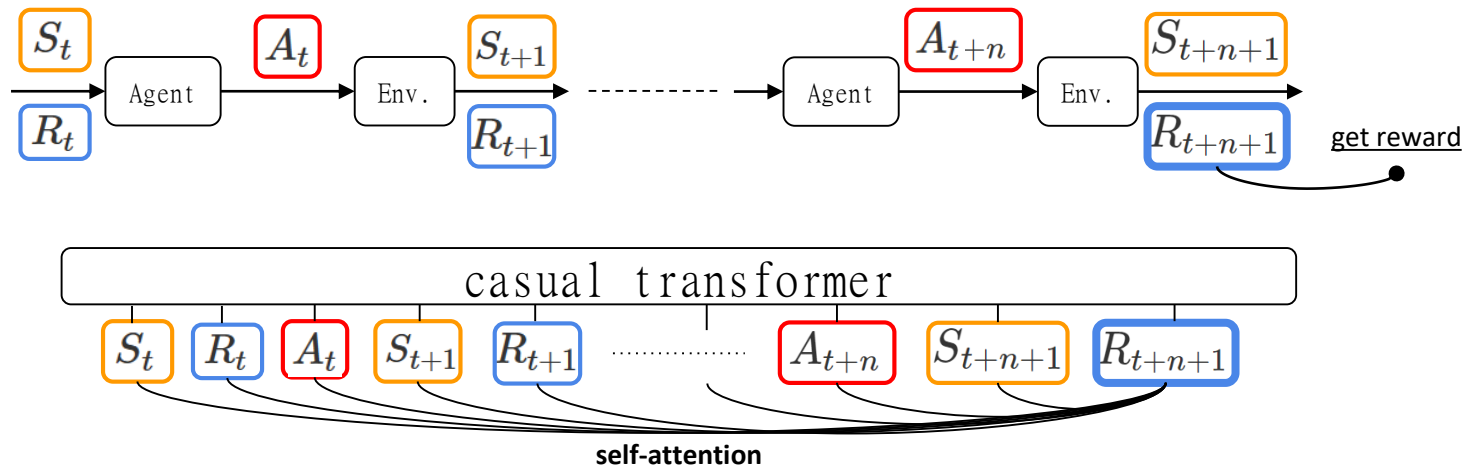
DoorKey task



2D Reacher (OpenAI Gym)

Motivation (sparse reward)

- Transformers can perform credit assignment directly via [self-attention](#), in contrast to Bellman backups which slowly propagate rewards and are prone to “distractor” signals.
 - This can enable transformers to still **work effectively in the presence of sparse or distracting rewards**.



Motivation: (short-sighted behaviors)

- Compared to when we solve RL problems under the **Model-free** condition in the past.
 - Decision Transformer allow us to **bypass the need for bootstrapping to propagate returns**
 - Avoids the need for **discounting future rewards**, as typically done in TD-learning, which can induce undesirable short-sighted behaviors.

Simplest temporal-difference learning algorithm:

TD(0)

– Update value $V(S_t)$ toward estimated return $R_{t+1} +$

$\gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

Update each value with the terminal result



Offline Reinforcement Learning

- Learning in a Markov decision process (MDP) described by the tuple (S, A, P, R).
 - S state, A actions, P state transition probability, R reward
- Instead of obtaining data via environment interactions
- The goal in reinforcement learning is to learn a policy which maximizes the **expected return** $\mathbb{E}\left[\sum_{t=1}^T r_t\right]$ in an MDP

Trajectory Representation

- The key desiderata in trajectory representation are :
 - should enable transformers to learn meaningful patterns
 - should be able to conditionally generate action at test time
 - generate actions based on **future desired returns**, rather than past rewards

$$\tau = \left(\hat{R}_1, s_1, a_1, \hat{R}_2, s_2, a_2, \dots, \hat{R}_T, s_T, a_T \right)$$

$$\hat{R}_t = \sum_{t'=t}^T r_{t'}$$

states $s \in \mathcal{S}$
actions $a \in \mathcal{A}$



How to define return-to-go?

1. We **specify** a target return based on our desired performance.
2. After executing the generated action, we decrement the target return by the achieved reward and obtain the next state.




Return-to-go conditioning

90 Breakout
2500 Qbert
20 Pong
1450 Seaquest

6000 HalfCheetah
3600 Hopper
5000 Walker
50 Reacher

\hat{R}_1

90



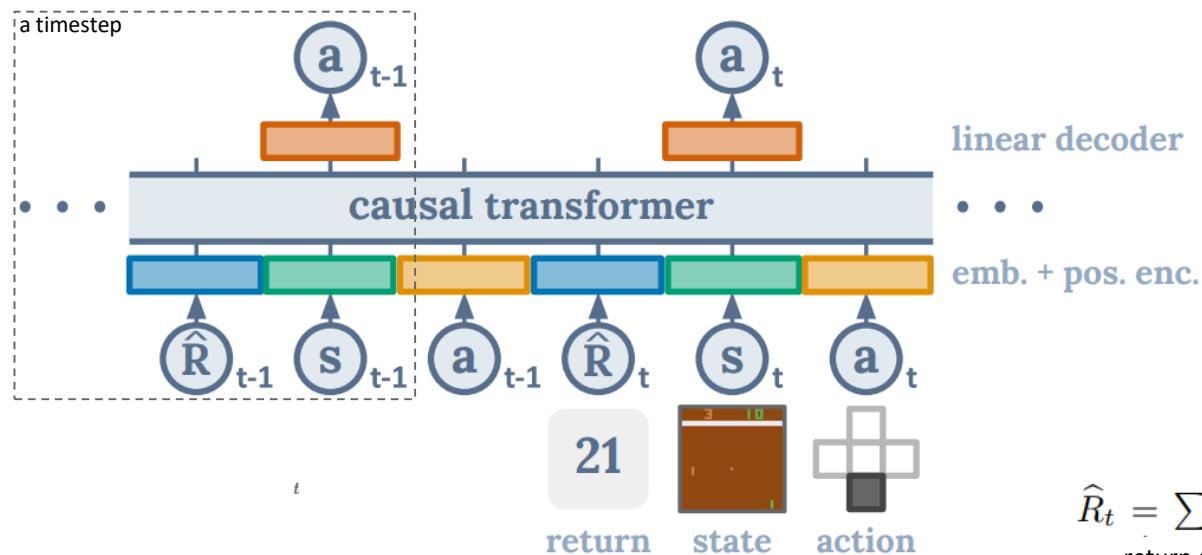
\hat{R}_2

85



Architecture

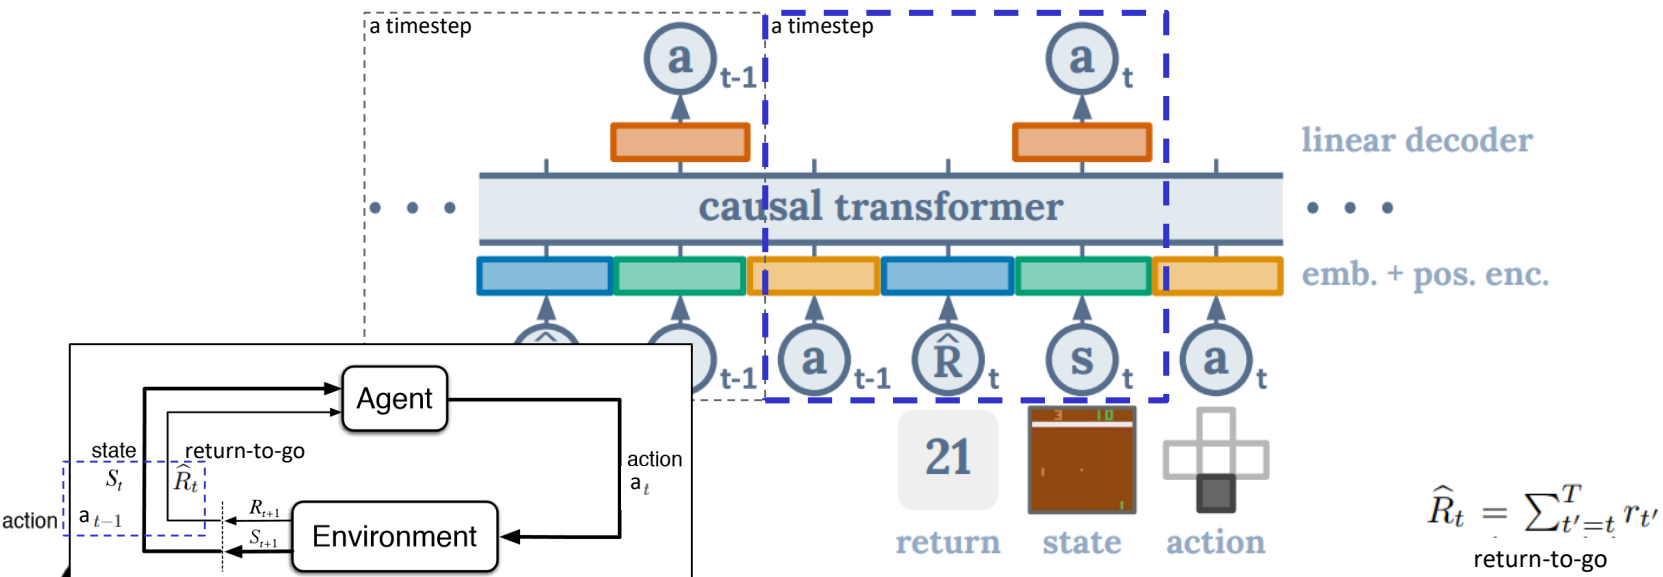
- Feed the last K timesteps into Decision Transformer, for a total of 3K tokens (for each modality include: return-to-go, state, action)



$$\hat{R}_t = \sum_{t'=t}^T r_{t'} \quad \text{return-to-go}$$

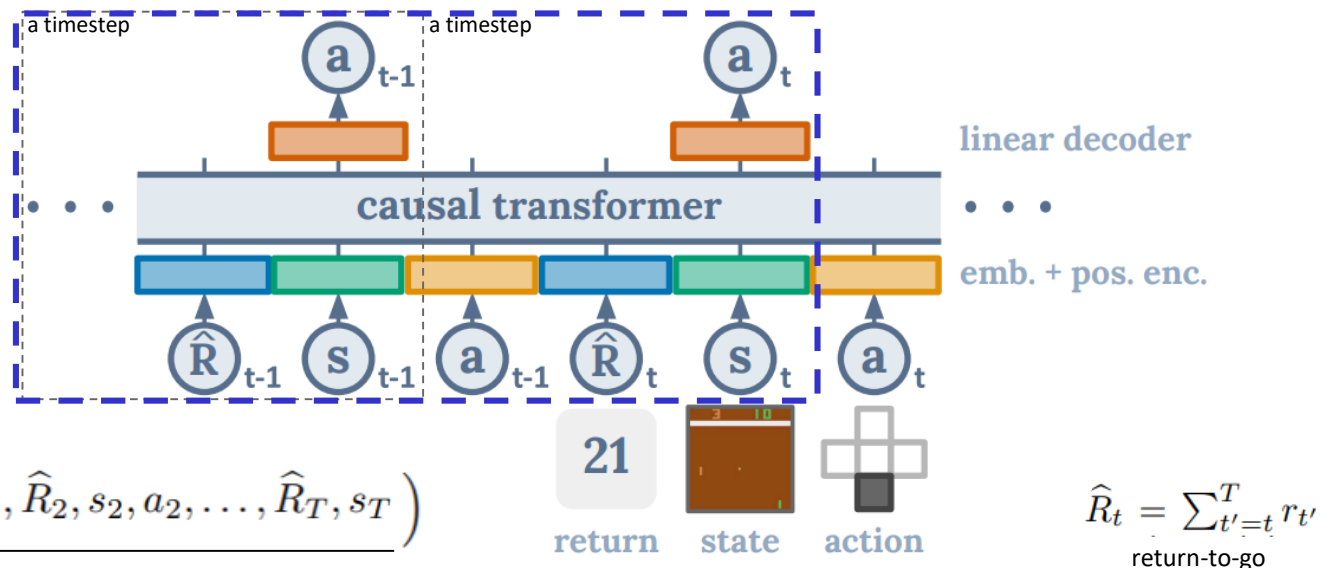
Architecture

- Feed the last K timesteps into Decision Transformer, for a total of 3K tokens (for each modality include: return-to-go, state, action)



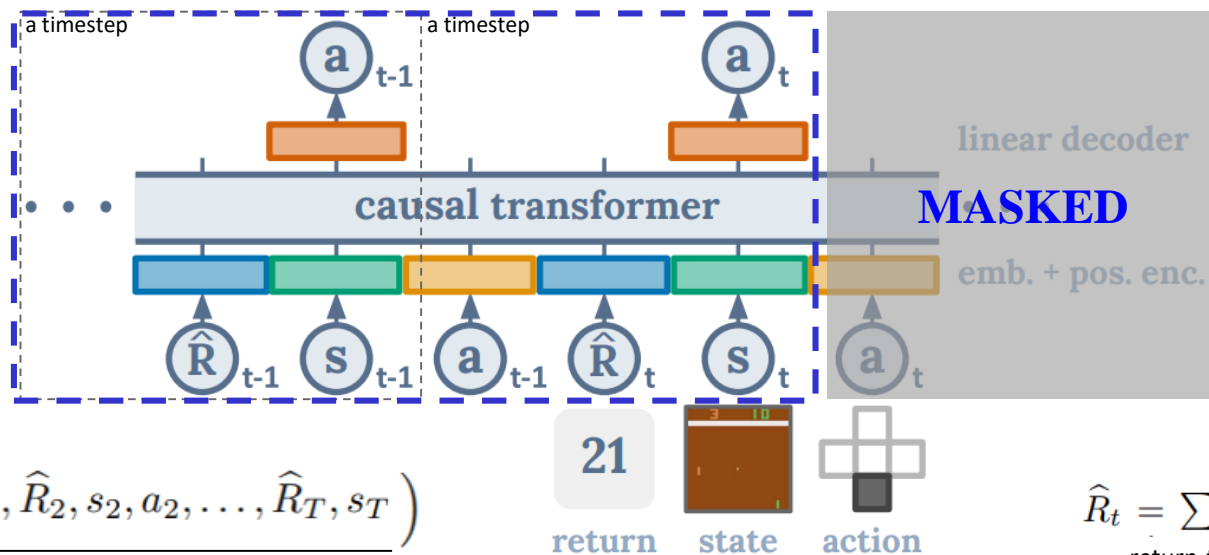
Architecture

- Feed the last K timesteps into Decision Transformer, for a total of 3K tokens (for each modality include: return-to-go, state, action)



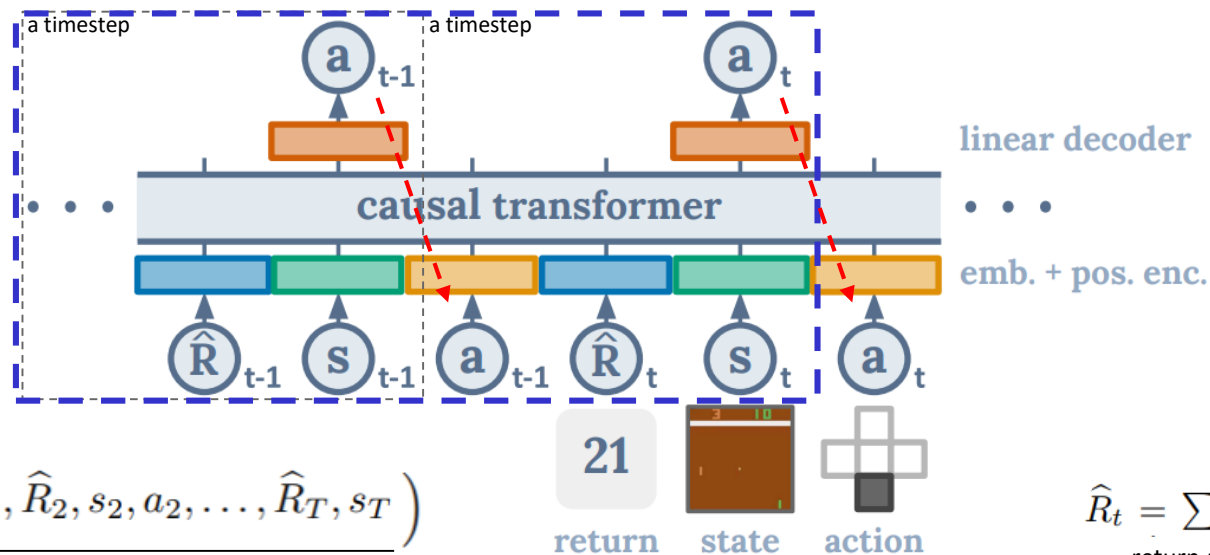
Architecture

- Feed the last K timesteps into Decision Transformer, for a total of 3K tokens (for each modality include: return-to-go, state, action)



Architecture

- Feed the last K timesteps into Decision Transformer, for a total of 3K tokens (for each modality include: return-to-go, state, action)



$$\hat{R}_t = \sum_{t'=t}^T r_{t'}$$

return-to-go



Experiments (Atari)

- Training Data: 1% of all samples in the DQN-replay dataset.
(500 thousand)



Context lengths of $K=30$ for Decision Transformer (expect $K=50$ for Pong)

Game	DT (Ours)	CQL	QR-DQN	REM	BC
Breakout	267.5 ± 97.5	211.1	21.1	32.1	138.9 ± 61.7
Qbert	15.1 ± 11.4	104.2	1.7	1.4	17.3 ± 14.7
Pong	106.1 ± 8.1	111.9	20.0	39.1	85.2 ± 20.0
Seaquest	2.4 ± 0.7	1.7	1.4	1.0	2.1 ± 0.3

BC : Behavior Cloning

CQL : Conservative Q-Learning

REM : Random Ensemble Mixture







Experiments (OpenAI Gym)

● Training Data: D4RL benchmark

Dataset	Environment	DT (Ours)	CQL	BEAR	BRAC-v	AWR	BC
Medium-Expert	HalfCheetah	86.8 \pm 1.3	62.4	53.4	41.9	52.7	59.9
Medium-Expert	Hopper	107.6 \pm 1.8	111.0	96.3	0.8	27.1	79.6
Medium-Expert	Walker	108.1 \pm 0.2	98.7	40.1	81.6	53.8	36.6
Medium-Expert	Reacher	89.1 \pm 1.3	30.6	-	-	-	73.3
Medium	HalfCheetah	42.6 \pm 0.1	44.4	41.7	46.3	37.4	43.1
Medium	Hopper	67.6 \pm 1.0	58.0	52.1	31.1	35.9	63.9
Medium	Walker	74.0 \pm 1.4	79.2	59.1	81.1	17.4	77.3
Medium	Reacher	51.2 \pm 3.4	26.0	-	-	-	48.9

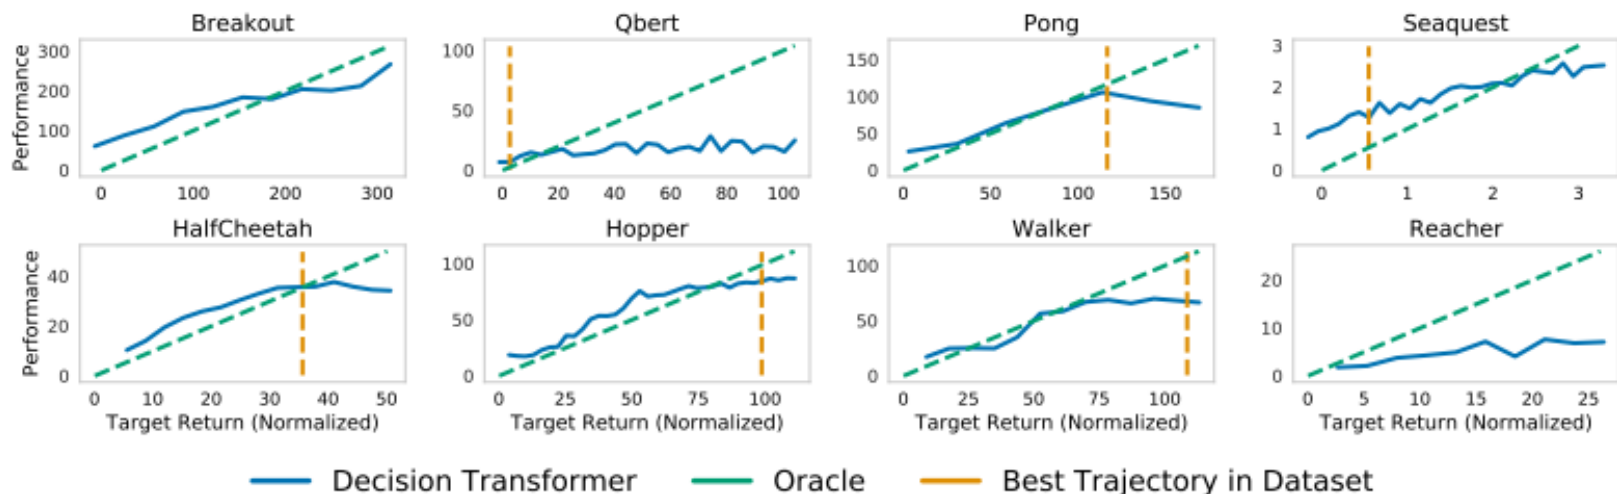
Medium-Expert:
1M: medium policy
1M: expert policy

Medium:
1M: medium policy



Experiments (Return-to-go)

- On every task, the **desired target returns** and the true observed returns are highly correlated.



Experiments (Behavior Cloning)

- Depends on environment, **Decision Transformer** can outperform %BC by using all trajectories in the dataset to improve generalization.
- Percentile Behavior Cloning (%BC) : Run BC on only the top X% of timesteps in the dataset.

Environment	DT (Ours)	10%BC	25%BC	40%BC	100%BC
HalfCheetah	42.6 \pm 0.1	42.9	43.0	43.1	43.1
Hopper	67.6 \pm 1.0	65.9	65.2	65.3	63.9
Walker	74.0 \pm 1.4	78.8	80.9	78.8	77.3
Reacher	51.2 \pm 3.4	51.0	48.9	58.2	58.4

Game	DT (Ours)	10%BC	25%BC	40%BC	100%BC
Breakout	267.5 \pm 97.5	28.5 \pm 8.2	73.5 \pm 6.4	108.2 \pm 67.5	138.9 \pm 61.7
Qbert	15.1 \pm 11.4	6.6 \pm 1.7	16.0 \pm 13.8	11.8 \pm 5.8	17.3 \pm 14.7
Pong	106.1 \pm 8.1	2.5 \pm 0.2	13.3 \pm 2.7	72.7 \pm 13.3	85.2 \pm 20.0
Seaquest	2.4 \pm 0.7	1.1 \pm 0.2	1.1 \pm 0.2	1.6 \pm 0.4	2.1 \pm 0.3



Conclusion

- Decision Transformer can match or outperform strong algorithms (CQL) designed explicitly for offline RL
- Minimal modifications from standard language modeling architectures